

MONGODB 3.6 DEVELOPER WORKSHOP

Palo Alto, CA – January 25th, 2018



SPONSOR ANNOUNCEMENTS





O'REILLY CONFERENCES

Strata Data Conf March 2018

AI Conference April 2018

Velocity & Fluent Conf June 2018

OSCON July 2018

JupyterCon August 2018

AI Conference September 2018

Strata


DATA CONFERENCE

March 6–8, 2018, San Jose, California

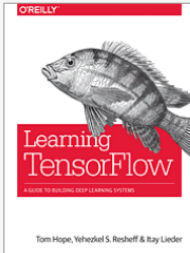
- All O'Reilly Conferences: <http://bit.ly/ORConfs>
- Use code **PCRAPHAEEL** to get a 20% discount on your ticket
 - available for ANY O'Reilly 2018 conference, not just Strata

FREE O'REILLY BOOK OFFER

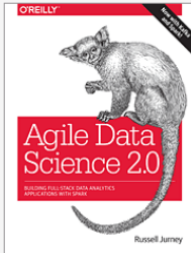
Choose one and download it at <http://bit.ly/mugbooks>



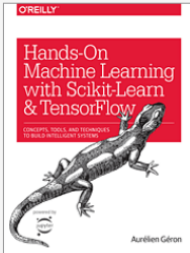
Data Science on the Google Cloud Platform
More Info
>



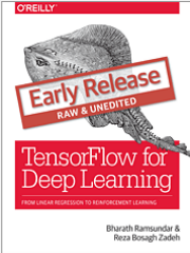
Learning TensorFlow
More Info
>




Agile Data Science 2.0
More Info
>



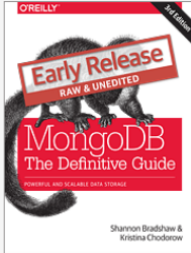
Hands-On Machine Learning with Scikit-Learn and TensorFlow
More Info >




TensorFlow for Deep Learning : Early Release
More Info
>



Designing Data-Intensive Applications
More Info
>



MongoDB: The Definitive Guide : Early Release
More Info
>



MongoDB Applied Design Patterns
More Info
>

MUST READS!

MONGODB EVENTS

Coming up soon in 2018...



#MDBlocal



MONGODB.local

MONGODB.LOCAL SEATTLE 02/15/2018

[HTTP://BIT.LY/MDBSEATTLE18](http://bit.ly/mdbseattle18)





MONGODB WORLD'18

BUILD GIANT IDEAS

New York City, June 27-28

Use code **RaphaelLondner** for a 25% discount
on the Super Duper Early Bird ticket (\$199)

www.mongodb.com/world Twitter hashtag: #MDBW18



WORKSHOP DETAILS





LOGISTICS

Free Wifi: xgen_public
Password: mongodbatlas

Bathroom is on your right after the door and your right (again) at the green phone booth.

YOUR WORKSHOP HOSTS




Raphael Londner

Developer Advocate
MongoDB

@rlondner

YOUR WORKSHOP HOSTS



Andrey Brindeyev

Technical Services Engineer
MongoDB

YOUR WORKSHOP HOSTS



Justin LaBreck

Senior Consulting Engineer
MongoDB




LET'S GET STARTED!

Workshop GitHub repository:

<http://bit.ly/mdb36workshop>

MongoDB Atlas Code: GOATLAS25

GETTING STARTED WITH MONGODB 3.6



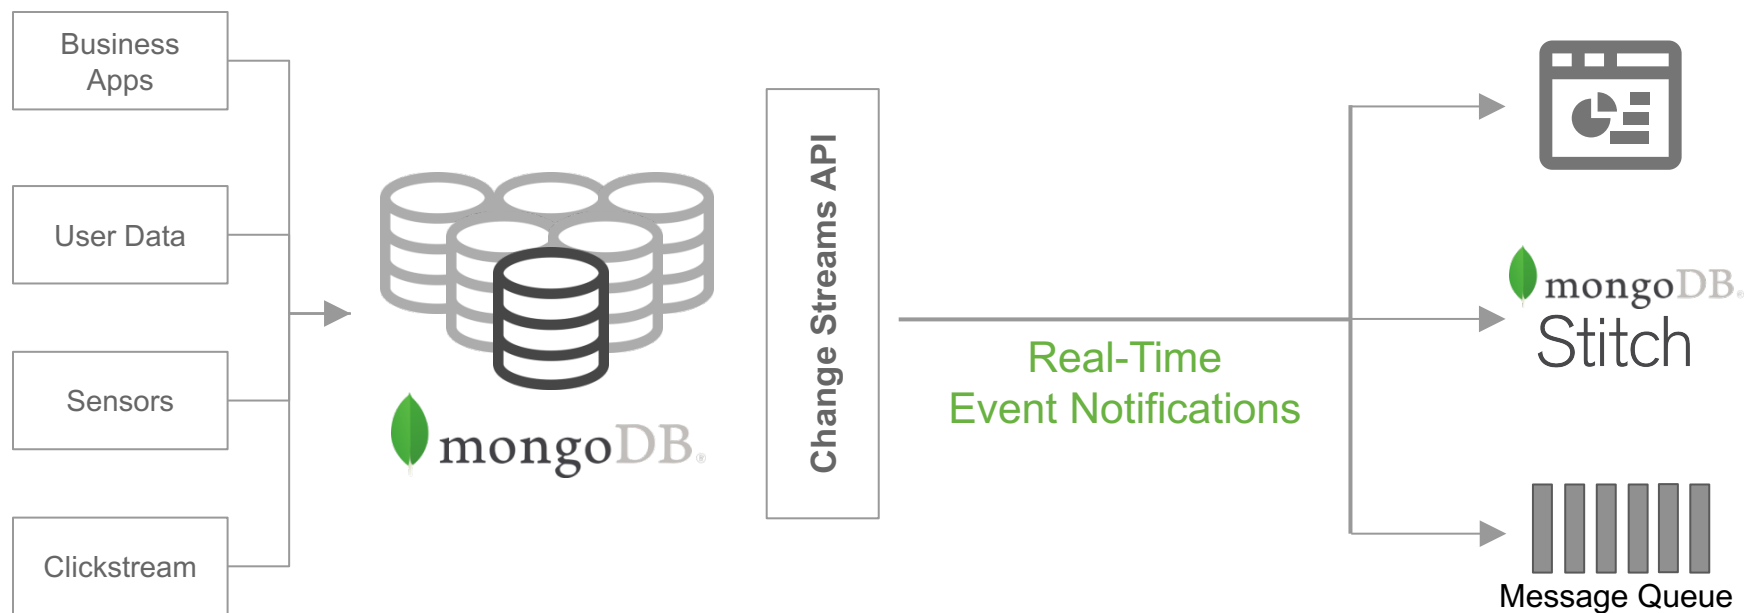
- [Download](#) the latest release & review the [Release Notes](#)
- Read the [What's New whitepaper](#)
- Free M036 [MongoDB University Training](#)
- We can help: [Major Version Upgrade Service](#)



Speed to Develop

- Change Streams
- JSON Schema Validation
- Query Expressivity
- Fully Expressive Array Updates
- Retryable Writes
- Tunable Consistency


MONGODB CHANGE STREAMS



Enabling developers to build
reactive, real-time services

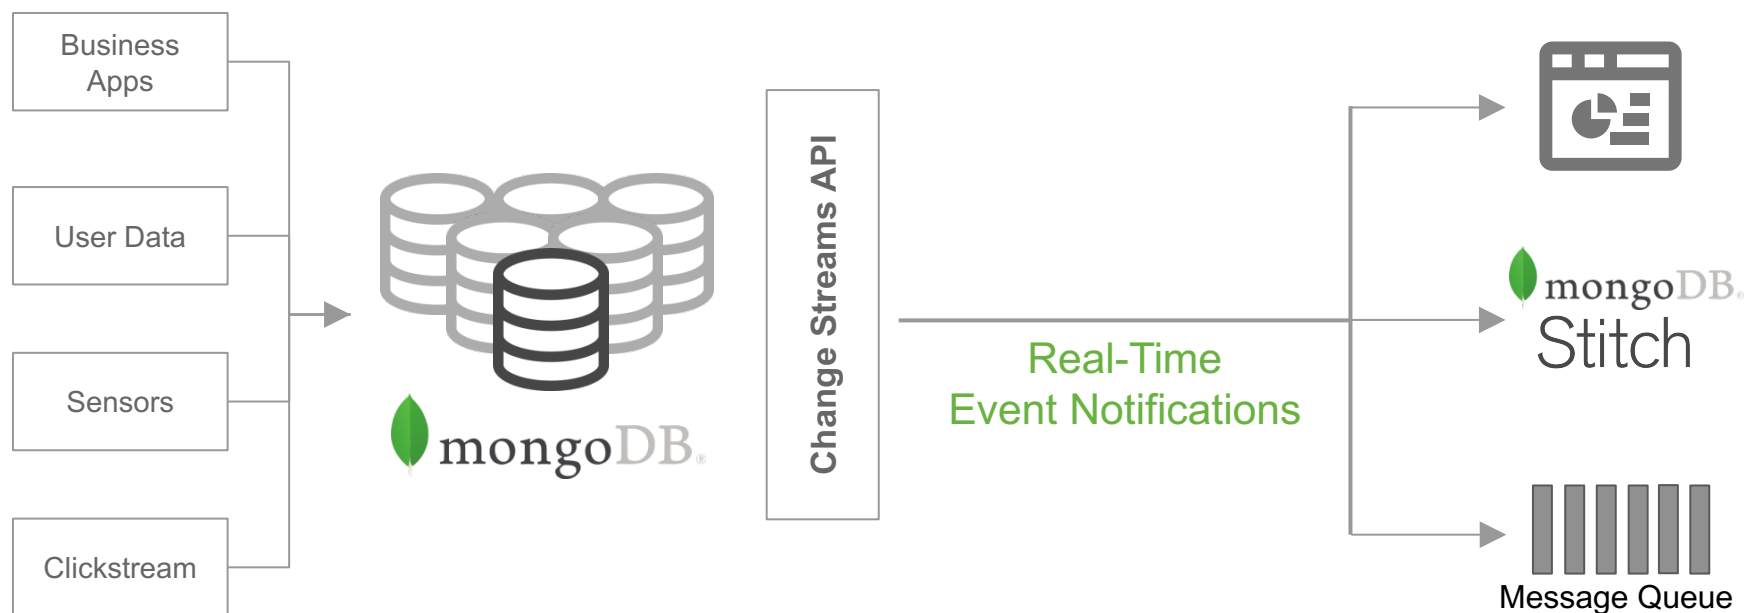
CHANGE STREAMS IMPLEMENTATION

Apps register for notifications via change streams API on top of MongoDB oplog



- Change streams are:
 - **Flexible:** deltas or the full document, filter on specific events only
 - **Consistent:** total ordering of data across shards
 - **Secure:** enforces collection's user access privileges
 - **Reliable:** only notifies once write committed on majority of replicas
 - **Resumable:** from node failure
 - **Concurrent:** up to 1,000 changes streams per MongoDB instance
 - **Familiar:** use regular MongoDB query language and drivers

CHANGE STREAMS USE CASES



- Refreshing trading apps as stock prices change
- Syncing changes across microservices
- Updating dashboards, analytics systems, search engines

- IoT data pipelines – e.g., generating alarms in response to connected asset failures
- Push new credit card transactions into ML models to recalculate risk
- Maintaining multiplayer game scoreboards



CHANGE STREAMS IN ACTION

```
// Select the collection to query.  
MongoCollection<Document> collection =  
    database.getCollection("orders");  
  
// Create the change stream cursor.  
MongoCursor<Document> cursor =  
    collection.watch().iterator();
```



SCHEMA VALIDATION IN ACTION

```
db.createCollection( "orders",
  {validator: {$jsonSchema:
    {
      properties:
        {line_items:
          {type: "array",
            items:
              {properties:
                {title: {type: "string"},
                 price: {type: "number", minimum: 0.0} },
                required: ["_id", "title", "price"],
                additionalProperties: false}}},
          required: ["line_items"]}}}}
)
```

SCHEMA VALIDATION

JSON Schema




Enforces strict schema structure over a complete collection for data governance & quality

- Builds on document validation introduced by restricting new content that can be added to a document
- Enforces presence, type, and values for document content, including nested array
- Simplifies application logic

Tunable: enforce document structure, log warnings, or allow complete schema flexibility

Queryable: identify all existing documents that do not comply

MONGODB RETRYABLE WRITES



Write failure handling moved from the app to the database for transient network errors or primary elections

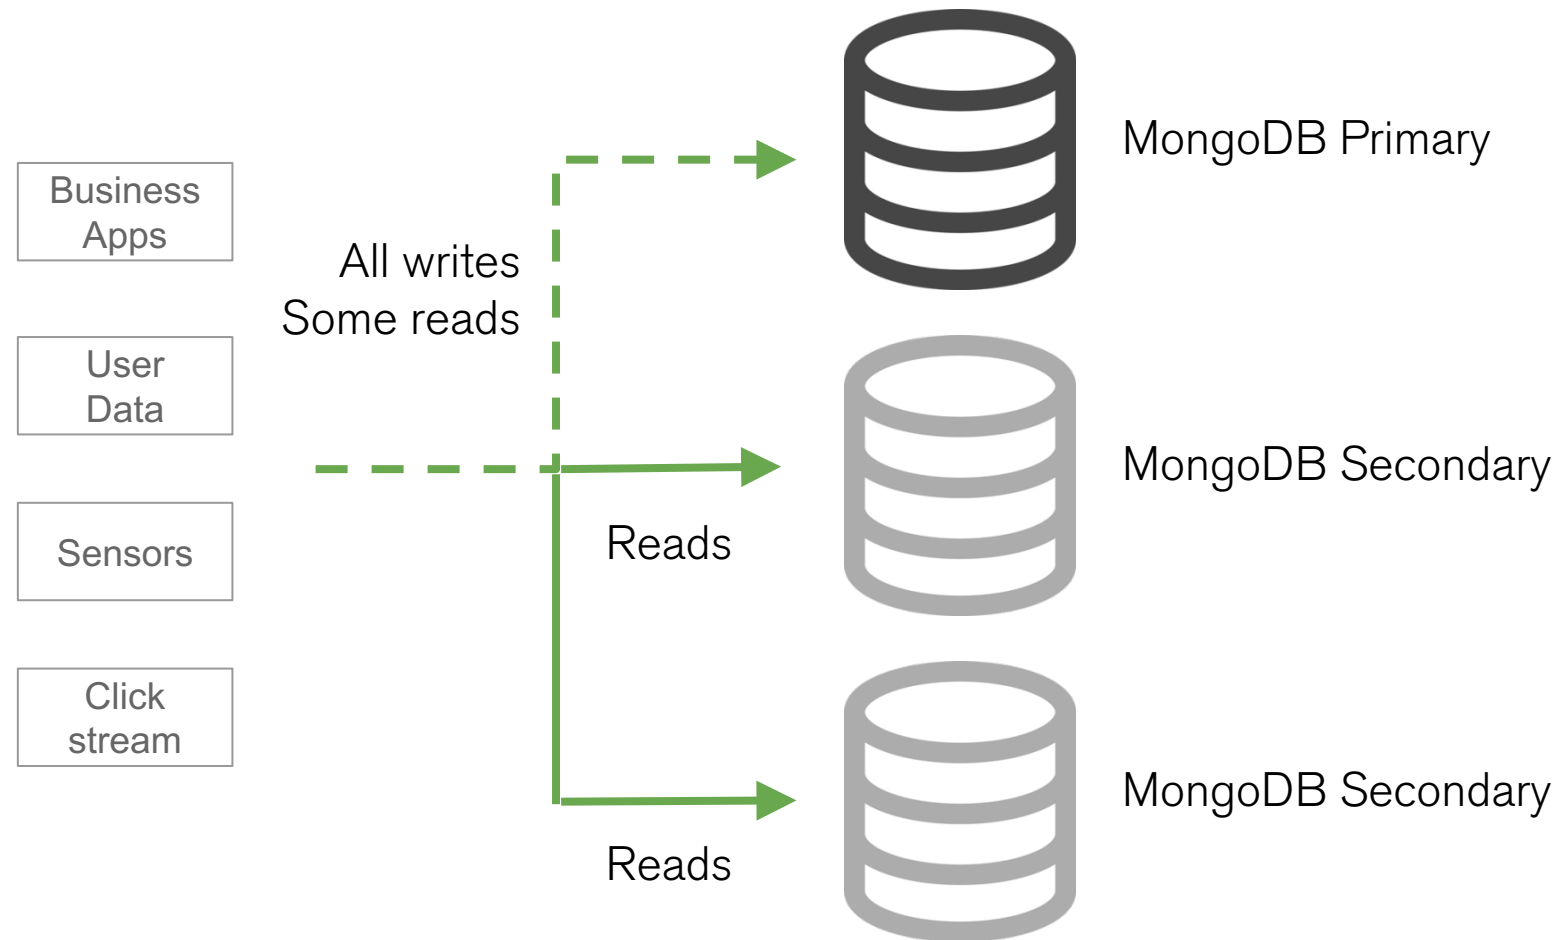
- Driver automatically retries failed write
- With a unique transaction identifier, server enforces exactly-once processing semantics
- Properties
 - Supports idempotent & non-idempotent operations, and errors caused by time-outs
- Delivers always-on, global availability of write operations
 - Overcomes the complexity imposed by multi-master, eventually consistent systems



RETRYABLE WRITES IN ACTION

```
uri = "mongodb://example.com:27017/?retryWrites=true"  
client = MongoClient(uri)  
database = client.database  
collection = database.collection
```

TUNABLE CONSISTENCY: SCALING READS



TUNABLE CONSISTENCY CONTROLS



Balance data consistency with performance SLAs

- Developers have precise control over how queries are routed across the database cluster
- **Causal consistency:** guarantees monotonic, logically consistent reads from any replica node in the same user session
- **Sharded secondary reads:** Secondary replicas now chunk-aware, ensuring consistent reads even as data is being rebalanced across a sharded cluster



CAUSAL CONSISTENCY IN ACTION

```
//start client session, which is causally consistent by default
try (ClientSession session =
    client.startSession(ClientSessionOptions.builder().build())) {

    //Run causally related operations within the session
    collection.insertOne(session, ... );
    collection.updateOne(session, ...);

    try (MongoCursor<Document> cursor =
        collection.find(session).filter(...).iterator()) {
        while (cursor.hasNext()) {
            Document cur = cursor.next();
        }
    }
}
```



QUERY EXPRESSIVITY & FULLY EXPRESSIVE ARRAY UPDATES

[]

- Use aggregation pipeline expressions within the MongoDB query language, using new `$expr` operator
 - SQL equivalent of `SELECT * FROM T1 WHERE a>b`
 - Example: find all customer accounts that have increased month on month spend by \$200 or more
 - More expressive queries with less client-side code
- Atomically update multiple matching elements of an array in a single update command
 - Example: update all prices in an array by 20%
 - More flexible data modeling
 - Avoids document rewrites imposed by other databases



UPDATING ARRAYS: ALL ELEMENTS

orders:

```
{
  _id: 5,
  line_items : [
    { id: 123,
      title : "USB Battery",
      price: 15.0 },
    { id: 512,
      title : "Hip T-shirt",
      price : 45.0 }
  ],
  ...
}
```

```
db.orders.update(
  { _id: 5 },
  { $mul: {
    "line_items.$[].price":
    0.8
  }
})
```




UPDATING ARRAYS: SOME ELEMENTS

orders:

```
{
  _id: 5,
  line_items : [
    { id: 123,
      title : "USB Battery",
      price: 15.0 ,
      shipped: true},
    { id: 512,
      title : "Hip T-shirt",
      price : 45.0,
      shipped: false }
  ]
}
```

```
db.orders.update(
  { _id: 5 },
  { $mul: {
    "line_items.$[li].price":
    .8}},
  {arrayFilters:[
    {"li.shipped":{$ne:true}}
  ]}
)
```

RICHER AGGREGATION PIPELINE



- **Expressive \$lookup**
 - Beyond Left Outer equi-join. Now supports non equi-joins & subqueries
 - Executed natively in the database, allowing more complex analytics queries with less code
- **Timezone-aware aggregations**
 - Enables multi-region analysis that are aware of region-specific timezones and working days when grouping data
- **New expressions for richer transformations**
 - Convert to and from objects to arrays of K-V pairs
 - Merge multiple objects into a single object
 - Remove fields from an object based on evaluation criteria

\$LOOKUP IN 3.6

orders:

```
{
  ...
  line_items : [
    { id: 123,
      title : "USB Battery",
      price: 15.0 },
    { id: 512,
      title : "Hip T-shirt",
      price : 45.0 }
  ],
  ...
}
```

```
db.orders.aggregate([
  {$unwind: "$line_items"},
  {$lookup:{
    from: "reviews",
    let: {p_id:
"$line_items.id"},
    pipeline: [
      {$match: {$expr: {$eq:
["$p_id", "$$p_id"]}}},
      {$group: {
        _id: 1,
        rating: {$avg: "$rating"}
      }}
    ],
    as: "avgRating" }
  }
])
```


MONGODB 3.6

MOVE AT THE SPEED OF YOUR DATA




Speed to Develop

- Change Streams
- Retryable Writes
- Tunable Consistency
- Compass
- Query Expressivity & Fine-Grained Array Updates



Speed to Scale

- Ops Manager
- Schema Validation
- Extended Security Controls
- E2E Compression
- Multi-Tenancy Management



Speed to Insight

- BI Connector
- Richer Aggregation Pipeline
- R Driver